

# **Towards Multi-Policy Hierarchical Scheduling in Linux for Containerized Space Applications**

Merlin Kooshmanian, Jérôme Ermont, Lucas Miné, Simon Corbin, Frédéric Boniol



RÉPUBLIQUE  
FRANÇAISE

*Liberté  
Égalité  
Fraternité*



# 1. Introduction

# 1. Introduction

---

*Why Using Linux for Space Applications ?*

## Current Solutions : RTOS and Hypervisors

- **Pros:** Deterministic, strong isolation, certifiable (DO-178C).
- **Cons:** Limited flexibility, restricted language/library support, static configuration.

## Linux

- **Pros:** Open-source, rich ecosystem (languages, libraries, tools), scalable, ...
- **Cons:** Limited isolation, unpredictable behavior for critical applications.

## Need

- Complement existing solutions by improving Linux guarantees:
  - Enhanced isolation.
  - Performance for heterogeneous workloads.



## 2. Context

# 2.1. Spacecraft Flight Software

## Diverse Functions

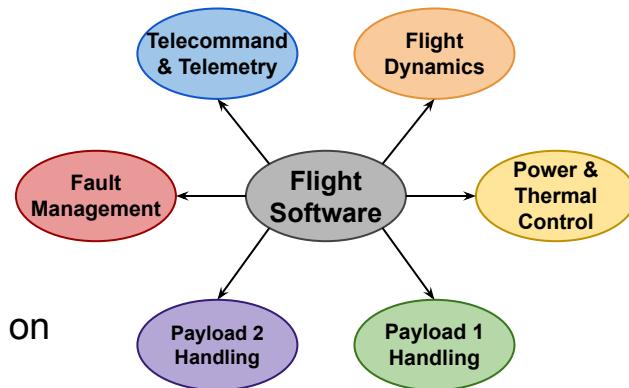
- Attitude and orbit control, power management, payload operations
- Monitoring, diagnostics, and fault management

## Heterogeneous Requirements

- Varying **criticality** levels
- Mixed **timing constraints**

## Needs

- **Isolation**
  - Prevent fault propagation by **separating** functions based on **criticality levels**.
  - Simplify development, validation, and integration by **reducing** interdependencies.
- **Optimized Scheduling**
  - Handle **heterogeneous workloads** with **mixed timing constraints**



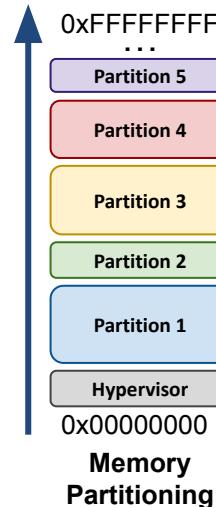
## 2.2. Isolation Concepts

### Spatial Isolation

Ensures that each component can **access only its explicitly assigned memory and hardware resources**, and **cannot read or modify the state of other components**.

#### Examples :

- Memory partitioning (static or dynamic)
- Access control & permissions

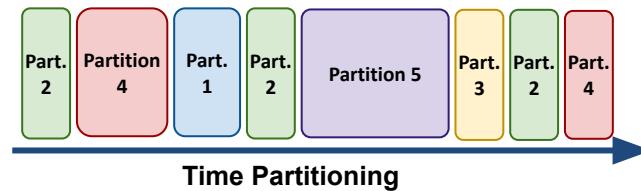


### Temporal Isolation

Ensures that the **timing behavior** of a component (e.g., execution time, response time, deadline satisfaction) **is independent** of the execution behavior **of other components**, within defined resource budgets.

#### Examples :

- Temporal partitioning
- Budget enforcement



*For our work, we adopt a **relaxed view of temporal isolation**, focusing on maintaining **bounded response-time guarantees** for tasks to ensure predictable timing behavior.*

## 2.3. Isolating Linux Applications

### Hypervisor-Based Isolation

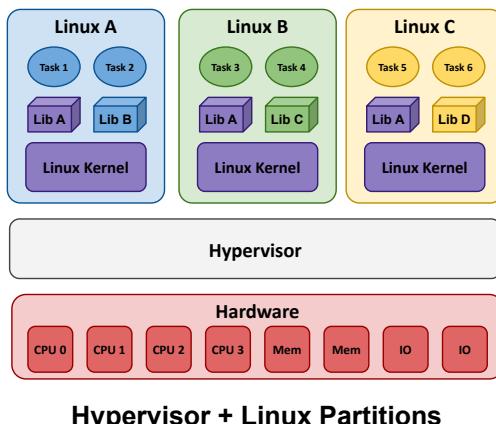
(Linux runs as a Guest OS in a virtualized partition)

#### Pros

- Strong spatial and temporal isolation

#### Cons

- Limited flexibility
- High memory footprint
- Hypervisor overhead



### Native Linux Isolation

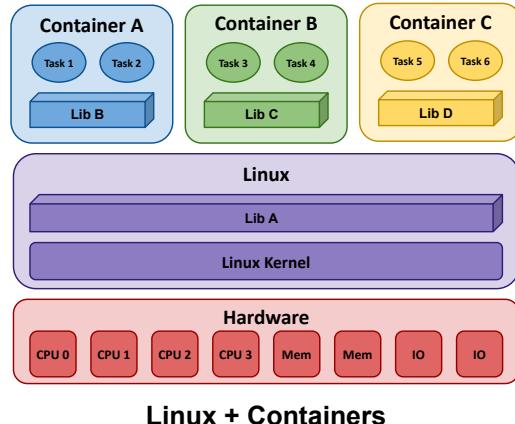
(Linux is the Host OS; isolation via processes, containers, etc.)

#### Pros

- Low memory footprint
- Less overhead

#### Cons

- Mostly software-based isolation
- Limited guarantees
- Weak temporal isolation



# 2.4. Linux Scheduling

## POSIX Scheduling Policies

- `SCHED_FIFO`, `SCHED_RR`, `SCHED_OTHER` defined by **POSIX**
- Linux extension: `SCHED_DEADLINE`

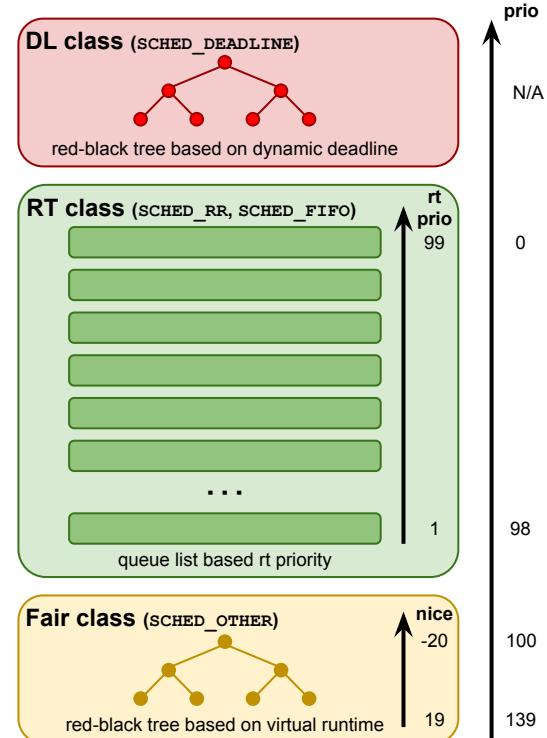
## Linux Scheduling Classes vs Policies

- **DL class** ⇔ `SCHED_DEADLINE` : EDF + Constant Bandwidth Server (CBS) for bandwidth enforcement
- **RT class** ⇔ `SCHED_FIFO`, `SCHED_RR` : fixed-priority preemptive scheduling (FIFO vs RR when same priority)
- **FAIR class** ⇔ `SCHED_OTHER` : fair CPU sharing

For convenience, policies will be grouped by this 3 classes

## Design

- Linux is a **multi-policy** OS : FAIR, RT, DEADLINE **coexist** at runtime
- Strict **class priority**: DEADLINE > RT > FAIR



## 2.5. Multi-Policy Scheduling

### Limits of Single-Policy Scheduling

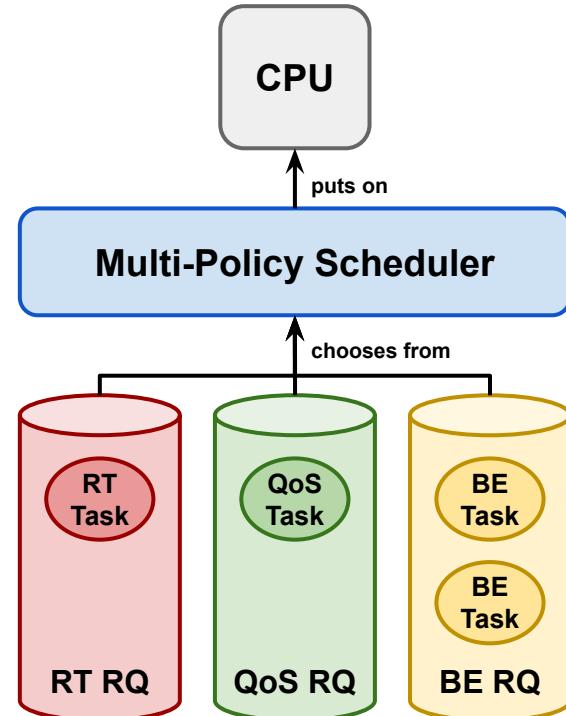
- Optimized for a **single objective** (e.g., reactivity, fairness, etc.)
- Inefficient for **heterogeneous software** needs

### Origins of Multi-Policy Scheduling

- Emerged in **general-purpose operating systems**
- Designed to meet **diverse user requirements**

### Key Benefits

- Provides **all scheduling policies** at runtime
- Enables **better fit** for **heterogeneous** workload requirements



## 2.6. Temporal isolation within Linux : HCBS patch

### Overview

- Hierarchical Constant Bandwidth Server (HCBS) enables hierarchical scheduling for the RT class
- Developed by **Abeni et al.** in 2019 [1] and improved by **Andriaccio et al.** in 2025 [2]

### Key Features

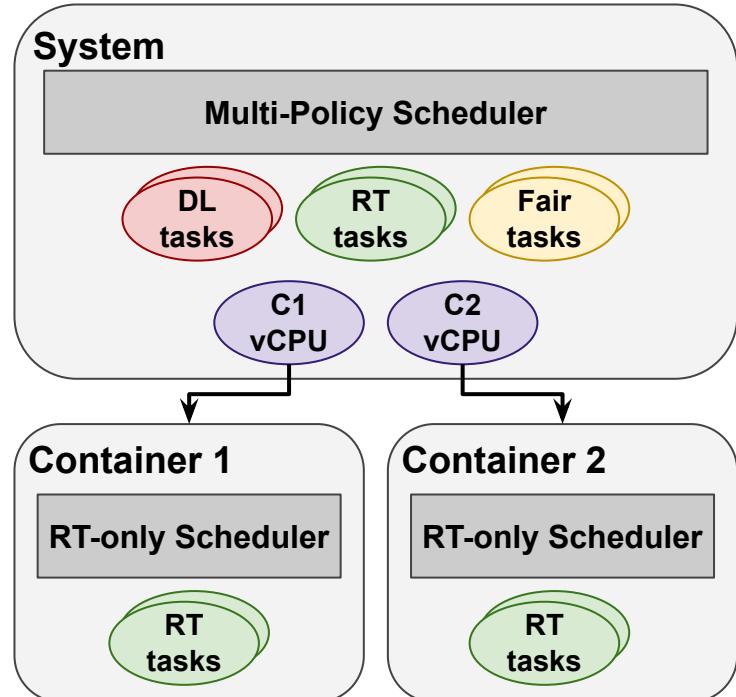
- Based on **deadline-servers** (CBS) to enable a **guaranteed bandwidth enforcement** between containers
- Associated **RT runqueues** with **deadline servers**, making it only suitable for RT tasks

### Limitation

- **Mono-policy**, relies solely on RT class
- **Not suited for task sets with heterogeneous objectives**

HCBS (original): <https://github.com/lucabe72/LinuxPatches/tree/HCBS>

HCBS (updated): <https://github.com/Yurand2000/HCBS-patch/tree/rt-cgroups>



→ HCBS creates a hierarchical scheduling in Linux

## 3. Problem Statement

### 3.1. Current Capabilities and Limitations with Linux

#### Linux-based Flight Software Needs

- **Spatial isolation**  
Strict resource separation (memory, CPU, devices)
- **Temporal isolation**  
Guaranteed CPU bandwidth and bounded latency
- **Scheduler for heterogeneous workloads**  
Ability to handle tasks with varying constraints

#### Native Linux Current Capabilities

- ✓ **Spatial isolation**  
Achieved via processes, cgroups, and namespaces
- ✗ **Temporal isolation**  
No guaranteed temporal isolation
- ✓ **Scheduler for heterogeneous workloads**  
Supports multiple scheduling policies

#### HCBS

- ✓ **Spatial isolation**  
Achieved via processes, cgroups, and namespaces
- ✓ **Temporal isolation**  
Hierarchical Scheduling including BW enforcement
- ✗ **Scheduler for heterogeneous workloads**  
Supports only RT tasks in containers

## 3.2. Objective: Hierarchical Multi-Policy Scheduling

---

### Problem Summary

Native Linux supports multi-policy scheduling but **lacks temporal isolation** guarantees. HCBS provides **hierarchical scheduling with bandwidth enforcement** but restricts scheduling to **real-time tasks only**.

### Objective

Extend hierarchical scheduling to support all Linux scheduling class (RT, DEADLINE, FAIR) within containers while **preserving**:

- **Temporal isolation** via bandwidth enforcement
- **Multi-policy flexibility** for heterogeneous workloads

## 4. Contribution

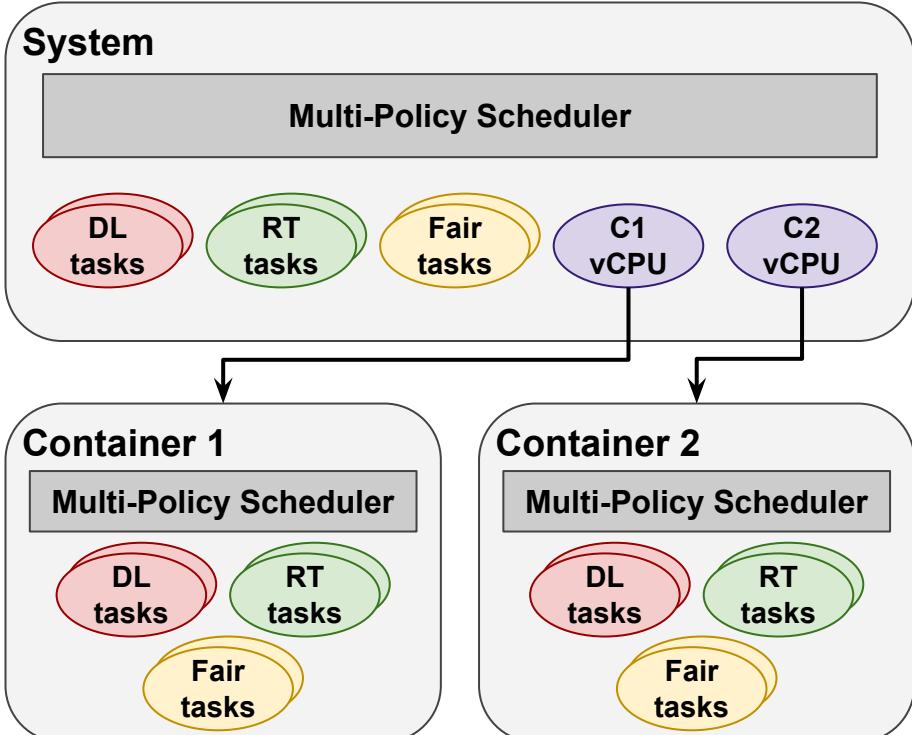
# 4.1. TGBS: Task Group Bandwidth Server

## Overview

- TGBS enforces container-level **temporal isolation** while **preserving multiple scheduling policy** within each container
- The core idea behind TGBS is to:
  - **Recreate scheduling** at the container level by **virtualizing CPUs**.
  - **Guarantee temporal isolation** between virtual CPUs.

## Limitation

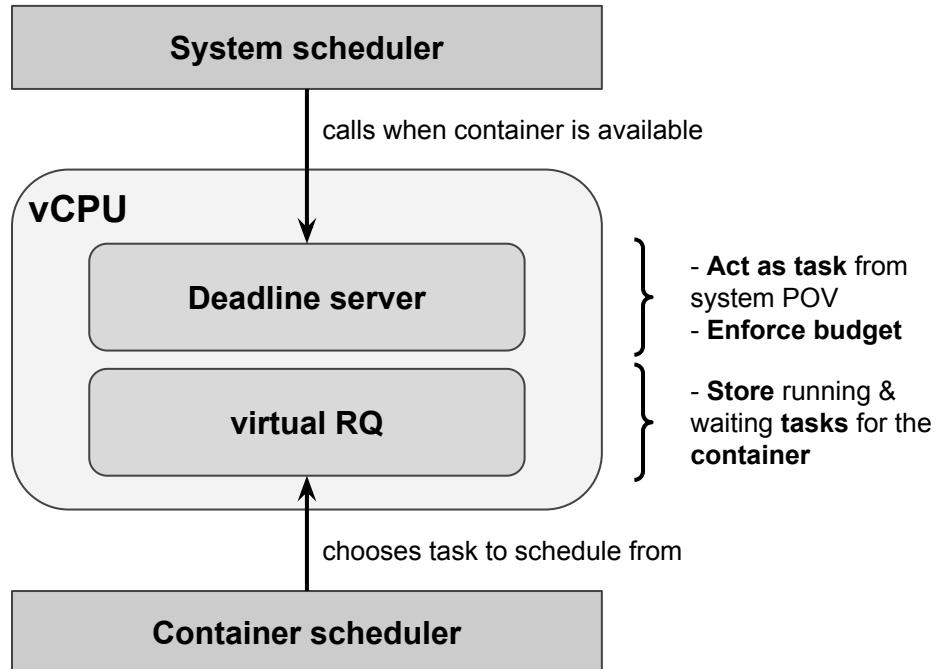
This work assumes a **SINGLE-CORE** environment. Multi-core scheduling is not yet supported and remains a direction for future research.



TGBS : <https://github.com/mkooshmanian/TGBS>

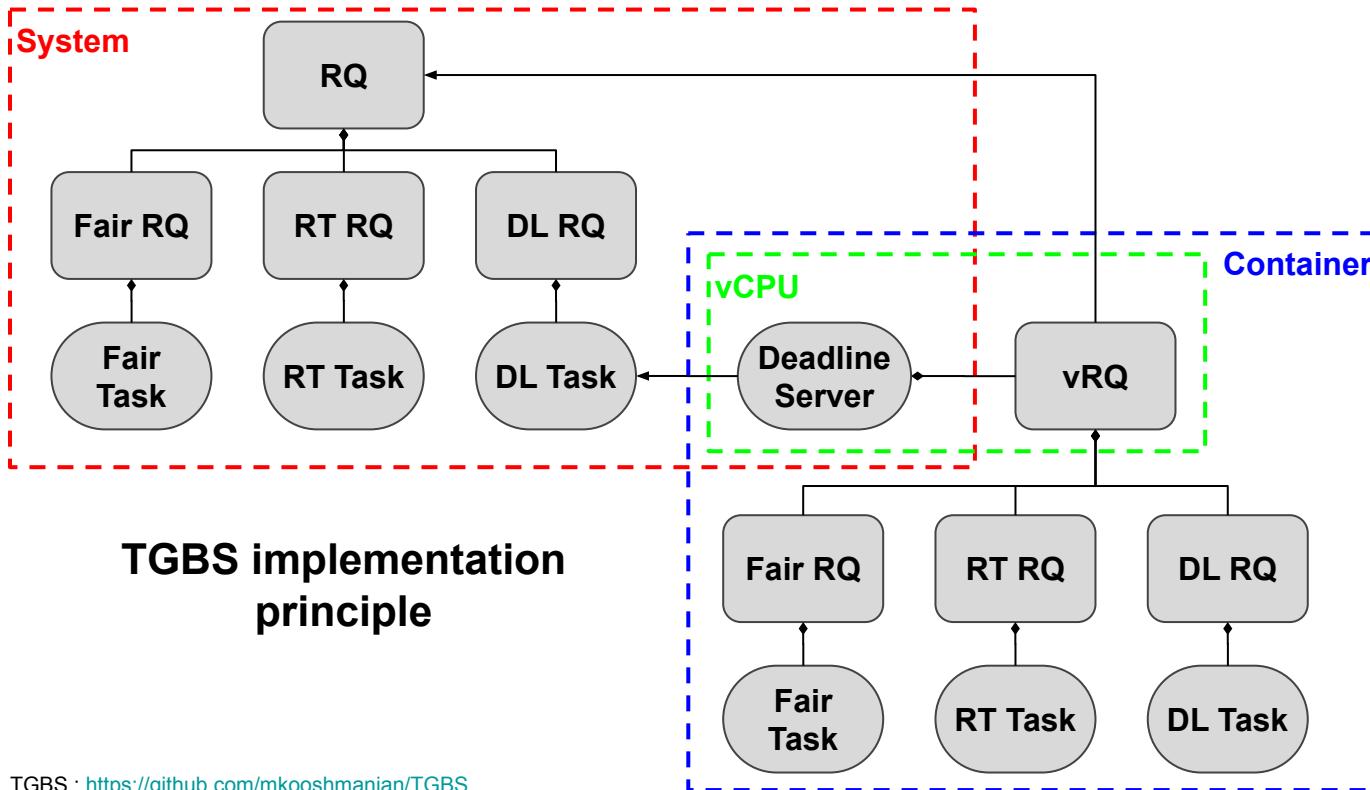
## 4.2. TGBS: CPU Virtualisation

- Rely on **deadline server** and **virtual runqueue**
- **deadline server**
  - is a system **scheduling entity**
  - **schedule by the Constant Bandwidth Server (CBS)** : BW enforcement
  - designed to **serve other tasks**
- **virtual runqueue**
  - **lightweight** version of a real runqueue



TGBS : <https://github.com/mkooshmanian/TGBS>

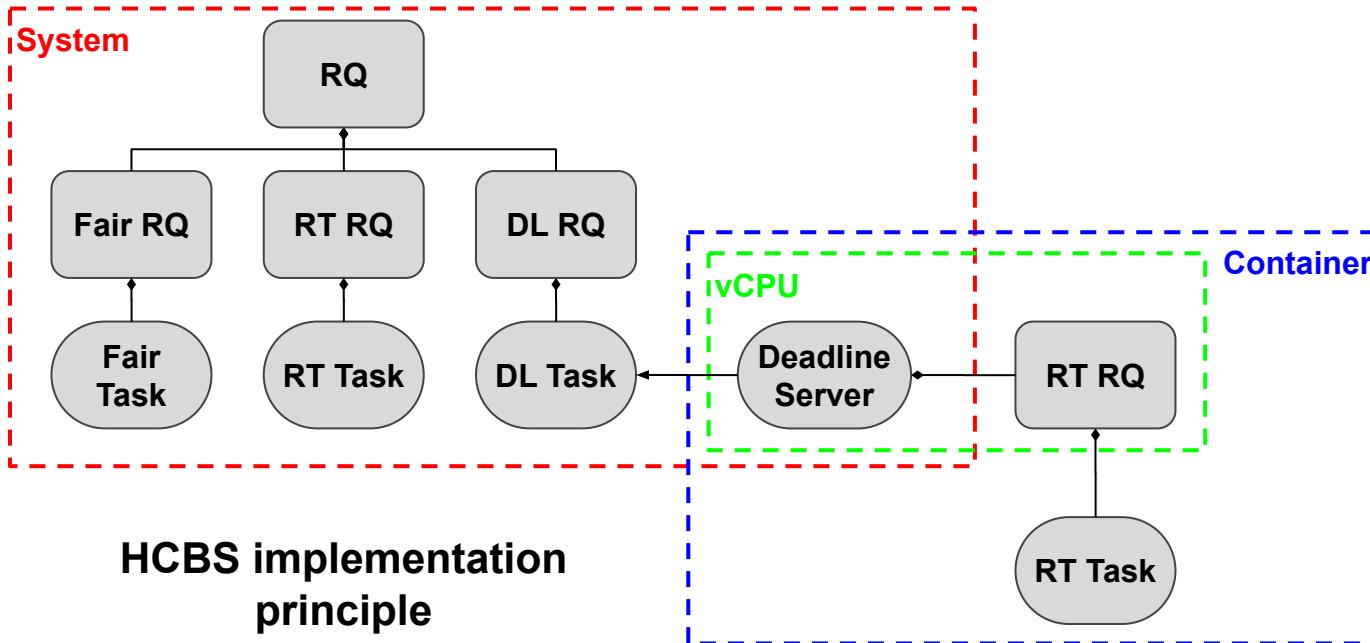
## 4.3. TGBS: Implementation Details



- TGBS **mirrors** Linux's **native scheduling structure** within containers.
- Virtual runqueues replicate class-specific queuing methods for all scheduling policies.

TGBS : <https://github.com/mkooshmanian/TGBS>

## 4.4. HCBS: Implementation Details (comparison)



- HCBS attaches **only an RT runqueue**, restricting it to real-time tasks.
- This **limits HCBS to mono-policy scheduling**.

HCBS (updated): <https://github.com/Yurand2000/HCBS-patch/tree/rt-cgroups>

## 5. Experimental Evaluation

# 5.1. Objective and Setup

## Objective

- Evaluate TGBS patch
  - Assess functional correctness and real-time performance
  - Compare with Vanilla Linux and HCBS

## Setup

- Zybo Z7 development board
  - Dual-core ARM Cortex-A9 (666 MHz)
  - 1 GB DDR3 memory
- Linux kernel v6.17 configured with PREEMPT-RT extension
- TGBS patch applied for hierarchical and multi-policy scheduling

disabled  
1 CPU



Zybo Z7 development board (Source: Digilent)

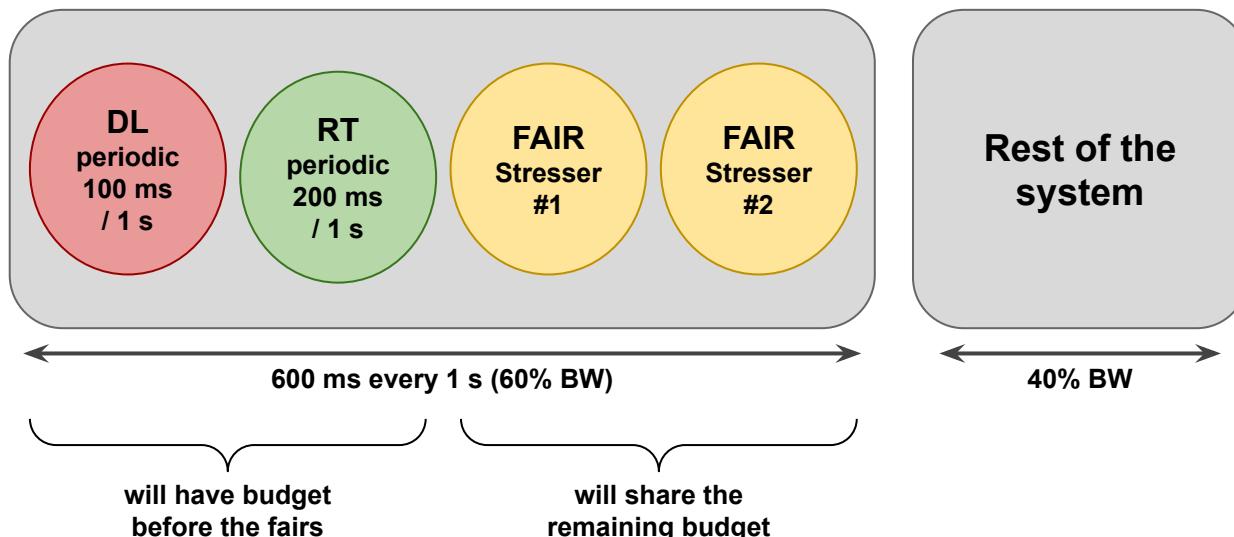
## 5.2. Functional Validation - Objectives

### Objective

- Verify correct prioritization between policies
- Ensure budget enforcement is working properly

### Setup

- 1 container with 4 threads
- 1 thread / policy → check multi-policy Stresser → check budget enforcement



### Expected results

- Prioritization**  
DL > RT > Fair
- Budget enforcement**  
 $U_{tot} \approx 60\%$
- Fairness**  
 $U(\text{stress#1}) \approx U(\text{stress#2})$

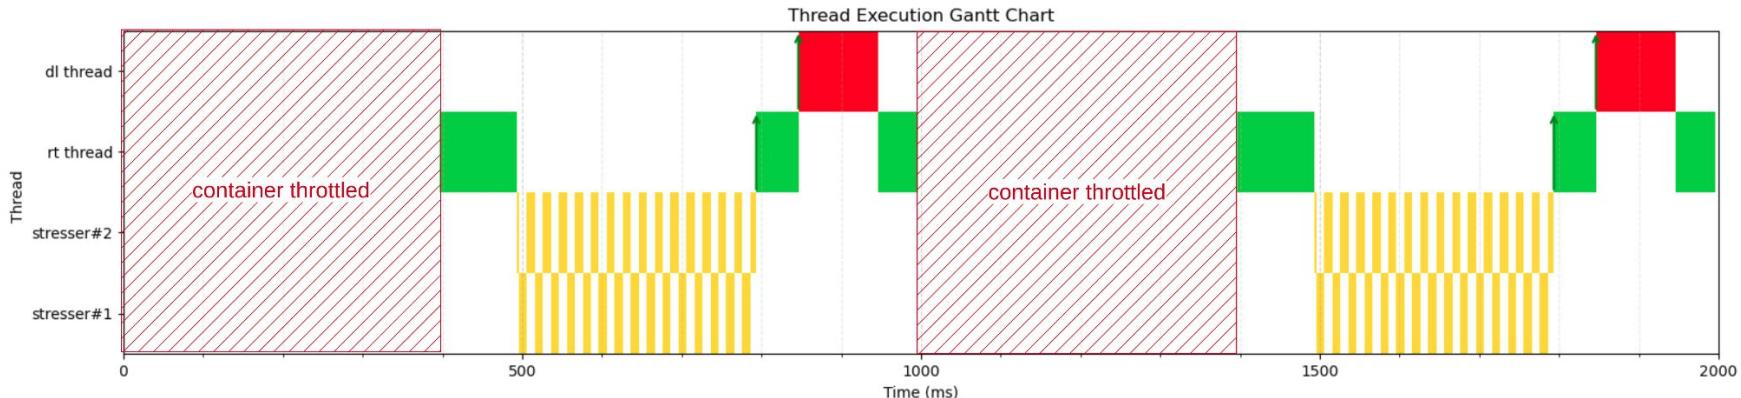
# 5.3. Functional Validation - Results

## Results

- Measured CPU shares match **expected proportions**
- Correct enforcement of **container reservation**
- Preemption between different **policies** is respected

CPU usage under a 600 ms / 1 s container reservation.

Task	Policy	CPU share
DL task (100ms / 1s)	DEADLINE	≈ 10%
RT task (200ms / 1s)	FIFO	≈ 20%
FAIR task (stresser #1)	OTHER	≈ 15%
FAIR task (stresser #2)	OTHER	≈ 15%



# 5.4. Overhead Evaluation - Objectives

## Objective

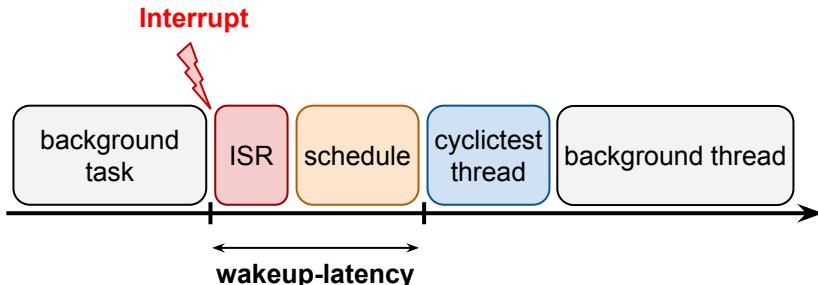
- Measure TGBS scheduling overhead
- Compare TGBS with Vanilla Linux and HCBS

## Methodology

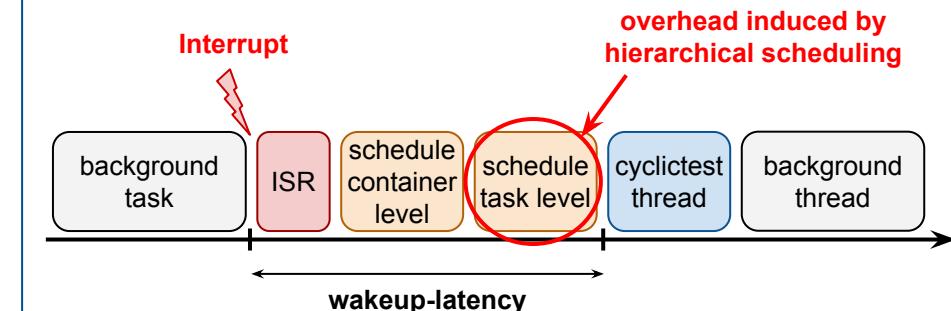
- Use **cyclictest** (rt-tests library) : evaluate **wake-up latency** (IRQ trigger to CPU execution)
- Run inside container for HCBS and TGBS

## Expectations

### Vanilla Linux



### TGBS or HCBS Linux



# 5.5. Overhead Evaluation - Results

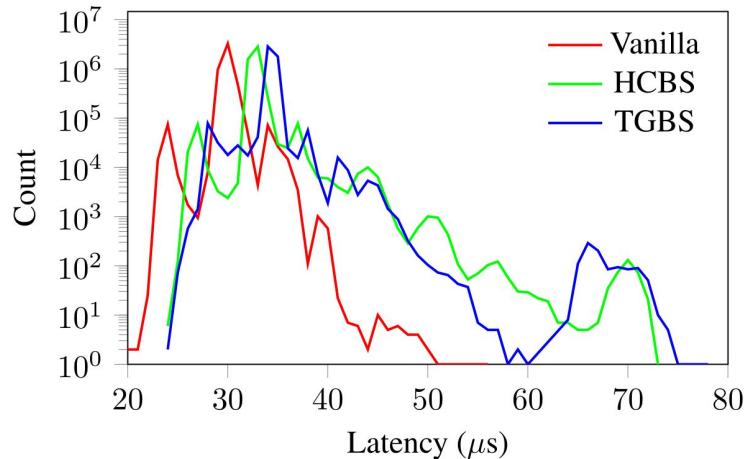
## Results

- Avg latency: 30  $\mu$ s (Vanilla)  $\rightarrow$  34  $\mu$ s (TGBS)  $\approx +15\%$
- Max latency: 56  $\mu$ s (Vanilla)  $\rightarrow$  78  $\mu$ s (TGBS)  $\approx +40\%$
- **TGBS overhead comparable to HCBS**
- **No pathological behavior** observed

Cyclictest Latency Statistics ( $\mu$ s)

System	Min	Avg	Max	Std Dev
Vanilla	20	30	56	1.3
HCBS	24	33	73	1.7
TGBS	24	34	78	1.5

Cyclictest Latency Distribution



## 6. Conclusion

# 6. Conclusion and Future Work

---

## Summary

- Demonstrated the **necessity** for **multi-policy** and **hierarchical scheduling** in future spacecraft flight software
- Developed and tested a Linux kernel patch (**TGBS**) to provide **temporal isolation** for **heterogeneous workloads**

## Future Work

- Address the **current gap** in literature regarding multi-policy and hierarchical models
- Develop a **formal model** of **TGBS** to prove its **correctness** and **timing guarantees**
- Extend TGBS to support **multicore** architectures

**Thanks for listening !  
Any questions ?**



*TGBS patch GitHub link*

contact : [merlin.kooshmanian@onera.fr](mailto:merlin.kooshmanian@onera.fr)



RÉPUBLIQUE  
FRANÇAISE

*Liberté  
Égalité  
Fraternité*



# Appendix

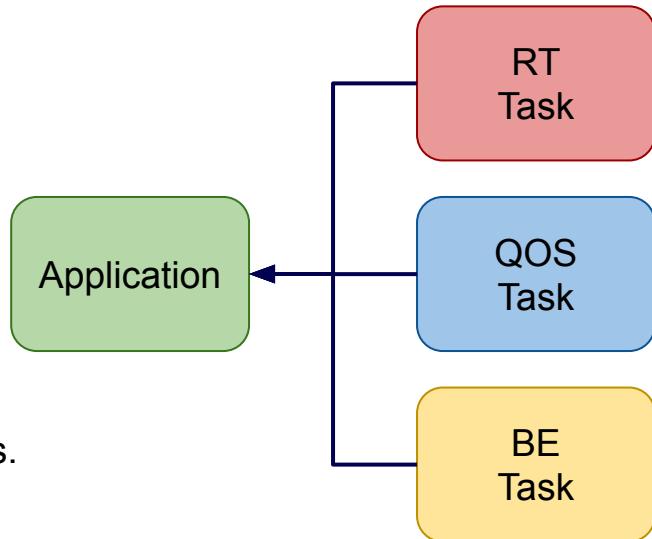
# A. Task Classes

## Diverse Operational Requirements

- **Heterogeneous** timing, reliability, and **operational demands** across space system functions
  - Predictability (critical control tasks)
  - Responsiveness (event driven tasks)
  - Fair sharing (variable workloads tasks)
  - Opportunistic execution (background tasks)

## Task Types

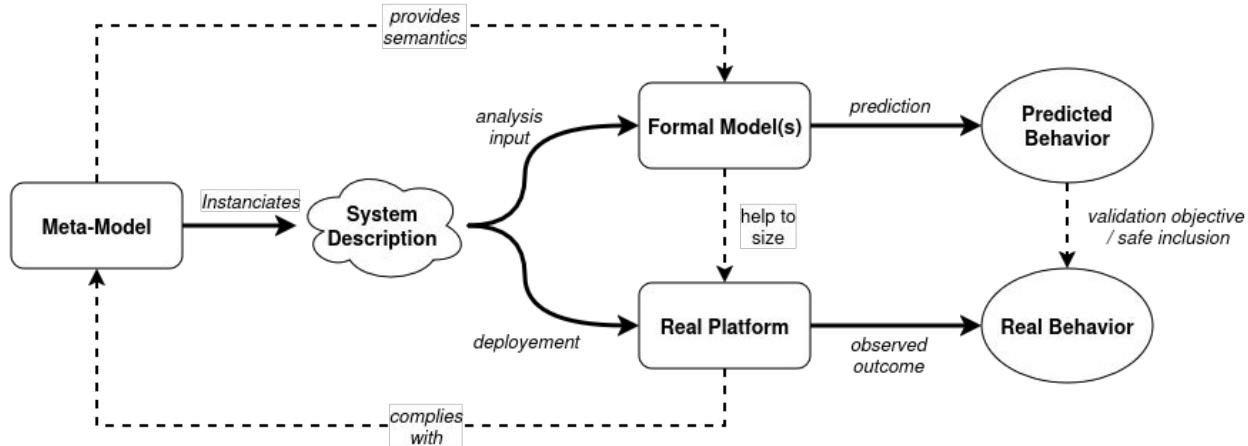
- **Real-Time (RT)**: Hard deadlines, zero tolerance for delays.
- **Best-Effort (BE)**: no needs, runs when possible
- **Quality-of-Service (QoS)**: Predictable progress, flexible timing.



# B. Methodology Overview

## Structured Process

- **Links conceptual scheduling systems to real platform execution**
- **Ensures coherence from concept definition to implementation**



## Core Elements

- **Meta-model / System Description**  
Defines and instantiates concepts
- **Real Platform**  
Executes the system description, yielding observed behavior

- **Formal Models :**
  - Provide conservative predictions for schedulability analysis
  - Provide real platform sizing methods

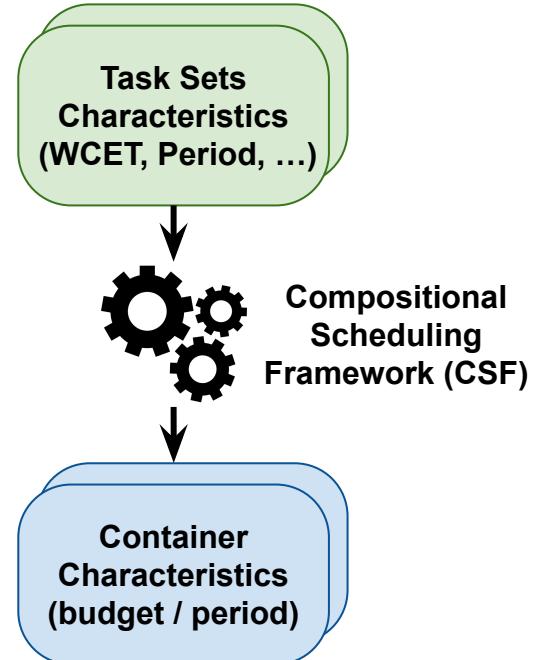
# C. Container Sizing – CSF

## Problem

- Need to **ensure** hard (RT) and soft (QoS) task **timing constraints** within container
- Guarantee **container schedulability** in **all situations** (migration, appearance, disappearance)

## Solution - Compositional Scheduling Framework

- Designed for hierarchical systems; **compositional**, ensuring container **schedulability independently**
- Based on **demand /supply mechanism**
- But **pessimistic** and assumes **mono-political scheduling** (each component has a mono-policy scheduler)



# D. Validation of the Container Sizing

## Objective

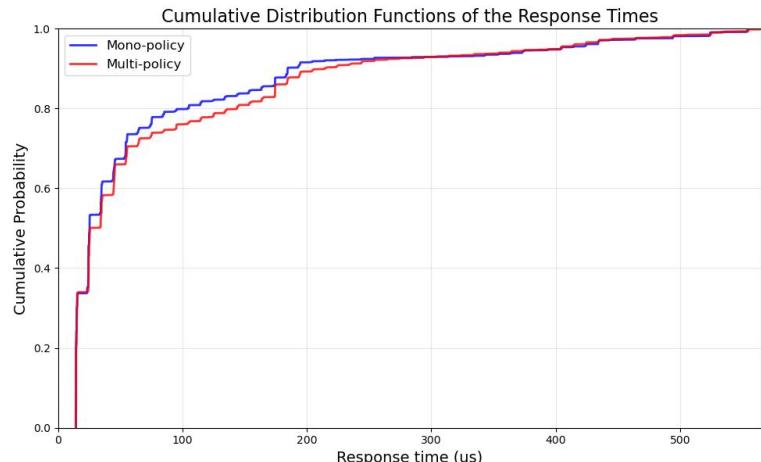
- Validate end-to-end methodology.
- Ensure correct execution under TGBS

## Methodology

- Synthetic workload of 10 periodic tasks, total utilization  $50\% \pm 5\%$
- Tasks distributed across two containers (A and B) sized with CSF assuming mono-policy.
- Mix of RT tasks (SCHED\_FIFO, RM priorities) and QoS tasks (SCHED\_OTHER)

## Results

- No deadline misses observed
- Analytically sized reservations sufficient
- Correct execution under both mono-policy and multi-policy configurations



Note : Mono-policy demonstrates slightly better response times than multi-policy, as expected since mono-policy prioritizes only system responsiveness, whereas multi-policy focuses also on fairness.